

MALLIKARJUN REDDY

GenAI / AI Engineer

hello@mallikarjunreddy.com | +91 8688715349 | [linkedin.com/in/infinity09](https://www.linkedin.com/in/infinity09) | github.com/infinity-009 | mallikarjunreddy.com

Summary

GenAI and Systems Engineer specializing in orchestrating Large Language Models (LLMs), deterministic state graphs via **LangGraph**, and scalable multi-modal AI architectures. Proven track record of architecture and end-to-end deployment of enterprise-grade RAG pipelines, fine-tuning state-of-the-art models (Claude, Llama) using LoRA/PEFT, and building highly-concurrent ML microservices. Adept at bridging bleeding-edge AI research with production-grade engineering using **Claude Code** for agentic pair programming to build programmatic systems.

Education

Indian Institute of Technology (IIT) Roorkee
Bachelor of Technology (CGPA: 7.4/10 | JEE Adv Rank: 3015)

Roorkee, India
2019 – 2023

Experience

Kredily

Senior GenAI & AI Systems Engineer

Jun. 2026 – Present

Bengaluru

- Built GenAI and RAG workflows for HR and payroll intelligence, enabling natural-language queries and automated reporting.
- Improved services handling roughly **20M requests/day**, reducing evening peak compute demand by about **40%** through caching and read-query optimization.
- Rolled out Attendance Regularisation V2 and the Payroll Arrears engine to all customers; strengthened Celery task and cloud workload reliability.

Isourse Technologies

AI/ML Engineer

Aug. 2024 – May 2026

Dwarka, New Delhi

- LLM Fine-tuning & Deployment:** Finetuned and deployed state-of-the-art open-source LLMs (Llama 3, Qwen, DeepSeek) using SFT, PEFT, and LoRA to architect a robust **Text-to-SQL conversational BI platform** (using **Pydantic**), completely automating enterprise dashboard generation.
- Agentic RAG Systems:** Developed a scalable **RAG-based AI chatbot** orchestrated with persistent agentic memory and vector stores (Milvus/pgvector), drastically improving contextual accuracy and driving a massive boost in user retention.
- AI-Powered ATS:** Engineered an **Intelligent ATS** system leveraging custom vector embeddings and semantic similarity, automating candidate evaluation and achieving near-human logical scoring accuracy at high latency efficiency.
- High-Volume Data Engineering:** Built and orchestrated a fault-tolerant, **real-time data warehousing ETL pipeline** (APIs → Kafka → PySpark) capable of seamlessly processing ultra-high-volume distributed data streams.
- MLOps & Infrastructure:** Spearheaded the migration of legacy ML infrastructure to highly concurrent, **Dockerized FastAPI** microservices, implementing cluster observability and alerting via Grafana and Prometheus.

Relevant Projects

ZENVideo (zenvidio.com) | *LangGraph, MCP, Claude, FastAPI, Next.js*

- End-to-End Multi-Agent Architecture:** Architected an autonomous video platform via **LangGraph** state graphs and **MCP** tool calls to generate Remotion code.
- Vision AI Feedback Loop & Scalable Render Engine:** Developed an "Agentic Mastery Loop" using **PIL/Vision AI** to auto-critique visual layouts and autonomously rewrite **.tsx** files, concurrently bridging Python AI backends with a Chrome/Puppeteer pipeline.

3D Avatar Motion Capture ([GitHub](https://github.com)) | *MediaPipe, Three.js, WebGL*

- Browser-Based Inference:** Architected an ultra-low latency, real-time 3D avatar system using **MediaPipe** for perfect alignment of full-body, hand, and complex facial mappings strictly via webcams.

Enterprise GenAI Chatbot | *Gemini 2.0, Llama 3.1, FastRTC, RAG*

- Multi-Modal Discovery:** Spearheaded a corporate AI assistant leveraging **hybrid search (semantic + exact match)** with Milvus and real-time bi-directional **FastRTC** voice capabilities to accelerate big-data extraction.

Technical Skills

Generative AI & ML: LLMs (Claude/GPT/Llama), LangGraph, Agentic Workflows, Text-to-SQL, Multi-Modal AI, LoRA, Vector Search (Milvus/pgvector)

Languages & Frameworks: Python, SQL, FastAPI, Pydantic, PyTorch, TensorFlow, LangChain, LlamaIndex, Node.js, React

Data & Cloud: Kafka, PySpark, PostgreSQL, Snowflake, Azure AI Foundry, GCP, AWS, Docker, Kubernetes

Developer Tools: Claude Code, Model Context Protocol (MCP), Git, CI/CD, ONNX, Grafana, Prometheus